

# SilentSpeak: A LipNet-style Visual Speech Recognition

Lipreading model, data preparation and training process (GRID corpus)

**Garvit (235890406)    Pranav (225890248)    Aadhisha  
Vaibhav**

October 28, 2025

## Abstract

SilentSpeak is a lipreading (visual speech recognition) project inspired by the LipNet architecture. The model combines 3D convolutional layers to capture spatiotemporal features from mouth-region video frames with stacked bidirectional LSTM layers and a CTC-based decoder to obtain character-level transcripts without frame-level alignment. This report details dataset preparation (GRID corpus), preprocessing, model architecture, training configuration, evaluation metrics and strategies used to mitigate speaker overfitting.

**Index Terms**— Visual Speech Recognition, Lipreading, 3D Convolutional Neural Networks, Bidirectional LSTM, Connectionist Temporal Classification, GRID corpus

## 1 Introduction

Visual speech recognition (lipreading) aims to translate visual mouth movements into textual speech. SilentSpeak implements a LipNet-style architecture that captures short-term spatiotemporal patterns using 3D convolutions and models longer-range temporal structure using bidirectional LSTMs. Training uses Connectionist Temporal Classification (CTC) to avoid explicit frame-level labeling. This report documents the dataset, preprocessing pipeline, model, training strategy and measures to reduce speaker overfitting.

## 2 Objectives

The main objectives of SilentSpeak are:

- Develop a lipreading pipeline that maps short video clips of the mouth region to character-level transcripts.
- Use a LipNet-style architecture (Conv3D encoder + BiLSTM decoder) and CTC loss for alignment-free training.
- Prepare a speaker-disjoint training/validation/test split on the GRID corpus to evaluate generalization.
- Document preprocessing, training hyperparameters and overfitting mitigation strategies.

## 3 System Overview

SilentSpeak comprises three logical modules:

1. **Preprocessing module:** extracts fixed-length frame sequences, crops to mouth region, converts to grayscale and normalizes.
2. **Model module:** 3D Conv encoder followed by BiLSTM layers and a time-distributed dense output producing per-frame logits.
3. **Training & evaluation:** CTC loss for training, greedy/beam CTC decoding for inference; metrics include WER and CER.

## 4 Dataset and Preprocessing

### 4.1 GRID corpus (reference)

The GRID corpus serves as the reference dataset for this project:

- **Speakers:** 34
- **Utterances per speaker:**  $\approx 1000$
- **Total duration:**  $\approx 34$  hours

### 4.2 Preprocessing steps

Each video sample is processed as follows:

1. Extract **75 frames** per utterance (3 seconds at 25 fps).
2. Crop to the mouth region (target size:  $140 \times 46$  pixels, width  $\times$  height).
3. Convert frames to grayscale (single channel).
4. Normalize pixel values to the range  $[0, 1]$ .
5. If a clip has fewer than 75 frames, pad by repeating the final frame; if longer, temporally sample or center-crop to 75 frames.

### 4.3 Annotations and splits

Each sample record contains `video_path`, `transcript` (character-level) and `speaker_id`. To prevent speaker memorization, create **speaker-disjoint** splits (e.g., 70% train, 15% validation, 15% test by speaker).

## 5 Model Architecture

SilentSpeak is built around three blocks: Conv3D encoder, Bidirectional LSTM stack, and time-distributed dense output for CTC.

### 5.1 Conv3D encoder

- Three Conv3D layers, each with **128 filters**.
- Kernel size:  $3 \times 3 \times 3$  (time, height, width).

- MaxPooling3D after each conv block (recommended pool size: (1, 2, 2) to preserve temporal resolution while reducing spatial dims).
- Batch normalization and SpatialDropout3D used for stability and regularization.

## 5.2 Temporal modeling (BiLSTM)

- Two stacked Bidirectional LSTM layers, **128 units** per direction.
- Return sequences = True (required for CTC).
- Dropout and recurrent dropout (typical 0.2–0.3) to reduce overfitting.

## 5.3 Output and decoder

- TimeDistributed Dense layer projects recurrent outputs to **vocab-size** logits per time-step (characters + blank symbol).
- Training uses **CTC loss**; inference can use greedy or beam-search CTC decoding.

## 5.4 Textual architecture diagram

```

Input (75 x 46 x 140 x 1) ->
[Conv3D(128) + MaxPool(1,2,2)] x3 ->
TimeDistributed(Flatten/Dense) ->
BiLSTM(128) -> BiLSTM(128) ->
TimeDistributed(Dense(vocab_size)) -> logits -> CTC

```

# 6 Training Configuration

## 6.1 Loss: CTC

Connectionist Temporal Classification (CTC) is used because it allows training without frame-to-label alignment. It sums probabilities over possible alignments and computes a negative log-likelihood.

## 6.2 Optimizer and hyperparameters

Hyperparameter	Typical value
Optimizer	Adam
Learning rate	$1 \times 10^{-4}$
Batch size	2–4 (GPU memory constrained)
Epochs	Up to 100 (with EarlyStopping)
Callbacks	ModelCheckpoint, ReduceLROnPlateau, EarlyStopping
Framework	TensorFlow / Keras
Hardware	NVIDIA GPU (e.g., GTX 1080 or better)
Estimated training time	~2–3 days (GPU, dataset dependent)

# 7 Evaluation Metrics

To quantify performance:

- **Character Error Rate (CER):** edit distance at character level normalized by reference length.
- **Word Error Rate (WER):** edit distance at word level normalized by reference word count.

## 8 Overfitting: Problem and Mitigation

### 8.1 Observed problem

A pre-trained LipNet-style model trained on some GRID speakers may overfit speaker-specific features (lip shape, lighting, viewpoint) and perform poorly on unseen speakers. This is the main practical challenge.

### 8.2 Mitigation strategies

- **Speaker-disjoint splits:** Ensure no speakers overlap between train/val/test.
- **Data augmentation:** temporal jittering, brightness/contrast variation, small spatial shifts, frame dropout/duplication.
- **Regularization:** SpatialDropout3D, L2 weight decay, dropout in LSTM layers.
- **Early stopping and LR scheduling:** Monitor WER/CER on validation speakers and use ReduceLROnPlateau/EarlyStopping.
- **Domain adaptation / adversarial training:** learn speaker-invariant features (advanced).
- **Speaker adaptation:** freeze encoder and fine-tune decoder on small speaker-specific data if available.

## 9 Extensions and Alternatives

- **Transformer encoder:** Replace BiLSTM with a Transformer to capture long-range dependencies using self-attention (requires more data/compute).
- **Attention-based seq2seq:** Use an encoder-decoder with attention rather than CTC (more flexible but more data hungry).
- **Multi-modal fusion:** Combine audio and visual streams for robustness when audio is available.

## 10 Practical Implementation Notes

- **CTC input lengths:** Ensure the `input_lengths` passed to CTC match the model's time-step outputs (adjust if time is downsampled).
- **Batch size:** Use small batches; use gradient accumulation if you need a larger effective batch size.

- **Checkpointing:** Save both training model (with CTC loss) and an inference model (produces logits).
- **Speaker evaluation:** Always evaluate on speaker-disjoint test set for realistic generalization numbers.

## 11 Conclusion

This report described the SilentSpeak LipNet-style visual speech recognition system: preprocessing for fixed-length mouth-region clips, a 3D convolutional encoder to capture spatiotemporal cues, stacked bidirectional LSTMs to model sequence dynamics, and a CTC-based training/decoding pipeline. Speaker overfitting is a primary concern when using GRID-like datasets; the recommended countermeasures include speaker-disjoint splits, augmentation, regularization, and careful validation using unseen speakers.